

Python + Prometheus = ?

Nikita Grishko



Flo Health Inc.

```
import prometheus
```

Monitoring?



Why Prometheus?  
Why not XYZ?

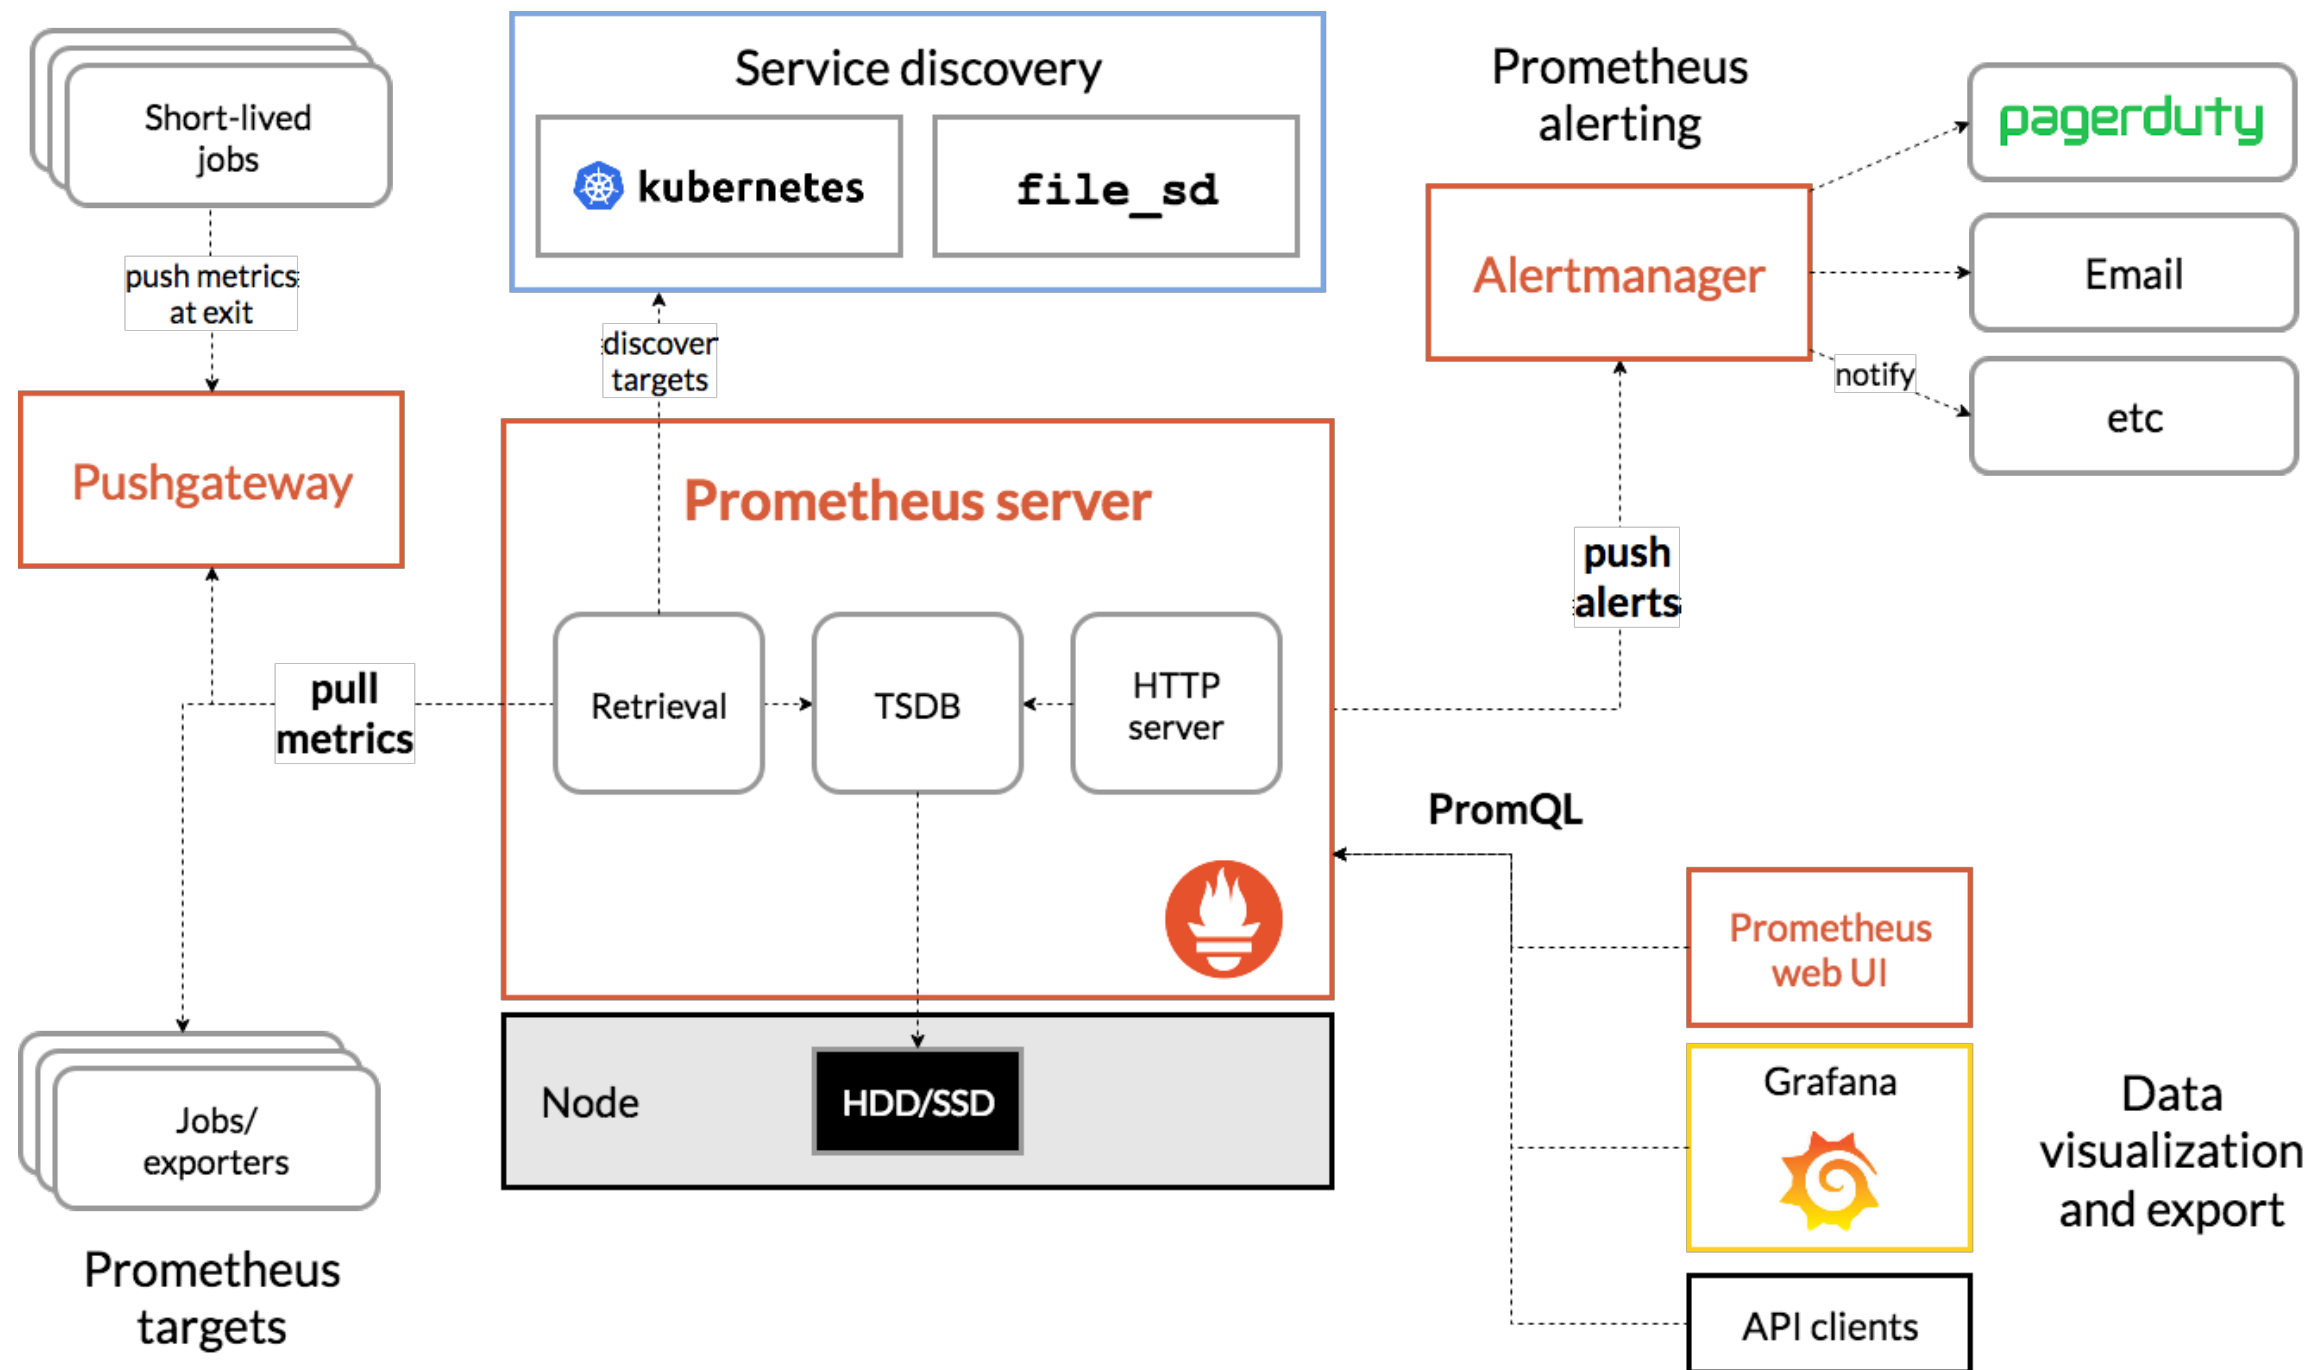
Two days !

Prometheus?

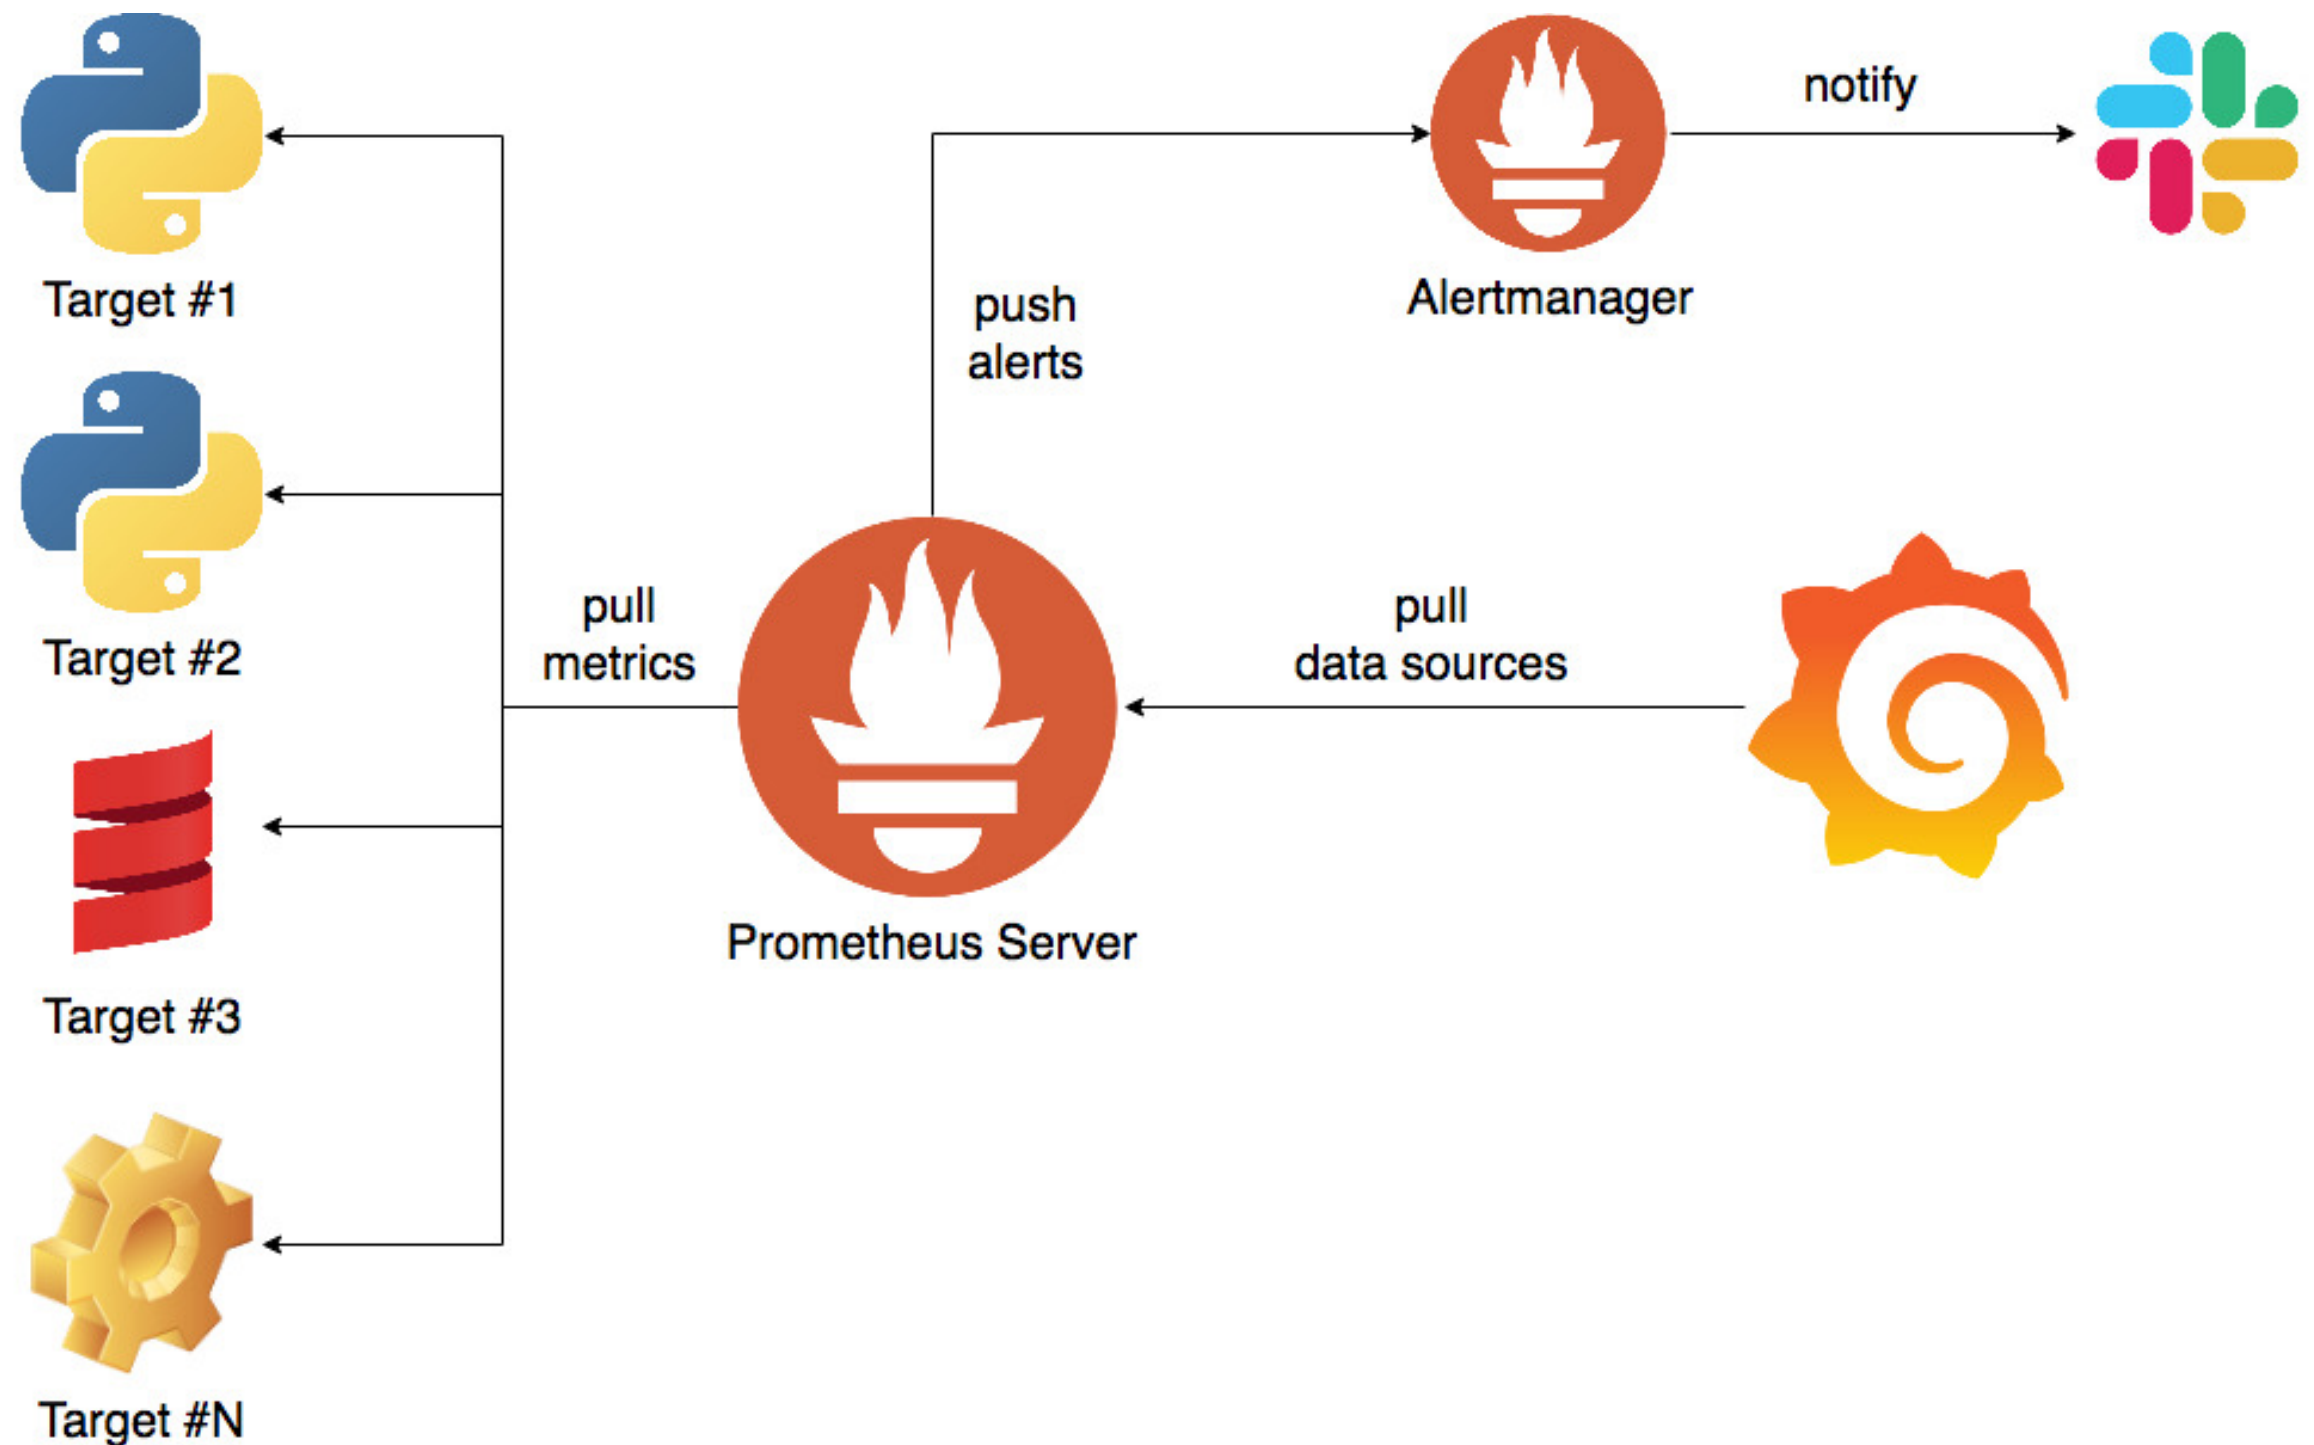


Prometheus is an open-  
source systems monitoring  
and alerting toolkit

# Prometheus Architecture



# Prometheus Architecture



# pull vs push

[https://prometheus.io/docs/introduction/faq/  
#why-do-you-pull-rather-than-push](https://prometheus.io/docs/introduction/faq/#why-do-you-pull-rather-than-push)

Let's code it!

```
1 from aiohttp import web
2 from prometheus_client import Counter, generate_latest
3
4
5 async def fire(request):
6     request.app["counter"].labels("fire").inc()
7     return web.Response(body=b"Fire!")
8
9
10 async def metrics(request):
11     request.app["counter"].labels("metrics").inc()
12     return web.Response(body=generate_latest())
13
14
15 async def get_application():
16     app = web.Application()
17     app.add_routes([
18         web.get("/fire", fire),
19         web.get("/metrics", metrics),
20     ])
21     app["counter"] = Counter("reqs", "Reqs counter", ["view"])
22     return app
23
```

```
$ python app_prom.py
```

```
$ for i in `seq 30`; do http :8080/fire; done
```



```
$ http :8080/metrics | grep reqs_total
# HELP reqs_total Req counter
# TYPE reqs_total counter
reqs_total{view="fire"} 30.0
reqs_total{view="metrics"} 1.0
```

```
$ python app_prom.py
```

```
$ gunicorn app_prom:get_application \  
    --workers=4 \  
    --worker-class=aiohttp.GunicornWebWorker
```

```
$ for i in `seq 30`; do http :8000/fire; done
```

```
$ http :8000/metrics | grep reqs_total
# HELP reqs_total Req counter
# TYPE reqs_total counter
reqs_total{view="fire"} 13.0
reqs_total{view="metrics"} 1.0
```

```
$ http :8000/metrics | grep reqs_total
# HELP reqs_total Req counter
# TYPE reqs_total counter
reqs_total{view="fire"} 9.0
reqs_total{view="metrics"} 1.0
```

Gunicorn is based on the pre-fork worker model. This means that there is a central master process that manages a set of worker processes.



***I'm going on an adventure!***



# Multiprocess Mode

- [https://github.com/prometheus/client\\_python#multiprocess-mode-gunicorn](https://github.com/prometheus/client_python#multiprocess-mode-gunicorn)
- The **prometheus\_multiproc\_dir** environment variable must be set to a directory that the client library can use for metrics.
- **MultiProcessCollector** must be used to collect files with metrics from file system **on each request** to /metrics endpoint.

# Multiprocess Mode

- Complex deployment configuration especially in case of containers...
- Multiprocessing mode is **slow**!

# Multiprocess Mode

- [https://github.com/prometheus/client\\_python/issues/374](https://github.com/prometheus/client_python/issues/374)
- [https://github.com/prometheus/client\\_python/issues/367](https://github.com/prometheus/client_python/issues/367)
- [https://github.com/prometheus/client\\_python/issues/275](https://github.com/prometheus/client_python/issues/275)
- [https://github.com/prometheus/client\\_python/issues/127](https://github.com/prometheus/client_python/issues/127)
- [https://github.com/prometheus/client\\_python/issues/204](https://github.com/prometheus/client_python/issues/204)
- <https://github.com/korfuri/django-prometheus/issues/89>
- ...



4GIFs.com

pull vs push

# Prometheus Pushgateway

- <https://github.com/prometheus/pushgateway>
- The Prometheus Pushgateway exists to allow ephemeral and batch jobs to expose their metrics to Prometheus.
- The Pushgateway is explicitly **not an aggregator or distributed counter** but rather a metrics cache.
- On each push, metrics **will be replaced**.

redis	← redis_exporter	← Prometheus
postgres	← postgres_exporter	← Prometheus
<b>xyz</b>	← <b>xyz_exporter</b>	← Prometheus

# Exporters

- `statsd + statsd_exporter`
- `collectd + collectd_exporter`
- `telegraf`



# What we want?

- Labels support.
- Easy to deploy and operate.
- Client libraries and community adaptation.

# collectd

- <https://github.com/collectd/collectd>
- [https://github.com/prometheus/collectd\\_exporter](https://github.com/prometheus/collectd_exporter)
- Limited tags (labels) support (<https://github.com/collectd/collectd/pull/1655>).
- Complex deployment (two services) and configuration.
- Abandoned client libraries.

# telegraf

- <https://github.com/influxdata/telegraf>
- [https://github.com/influxdata/telegraf/tree/master/plugins/outputs/prometheus\\_client](https://github.com/influxdata/telegraf/tree/master/plugins/outputs/prometheus_client)
- Hard to configure aggregations and integration with prometheus.

# statsd


- <https://github.com/statsd/statsd>
- [https://github.com/prometheus/statsd\\_exporter](https://github.com/prometheus/statsd_exporter)
- Labels not supported.
- Complex deployment (two services) and configuration.

# statsd\_exporter

- Since the StatsD exporter uses the same line protocol as StatsD itself, you can also configure your applications to send StatsD metrics directly to the exporter. In that case, **you don't need to run a StatsD server anymore.**

# statsd\_exporter

- DogStatsD extension (<https://docs.datadoghq.com/developers/dogstatsd/>).
- `name:value|type|@sample_rate|#tag:value`
- The exporter will convert DogStatsD-style tags to prometheus labels.
- <https://pypi.org/project/datadog/>



Let's try it!



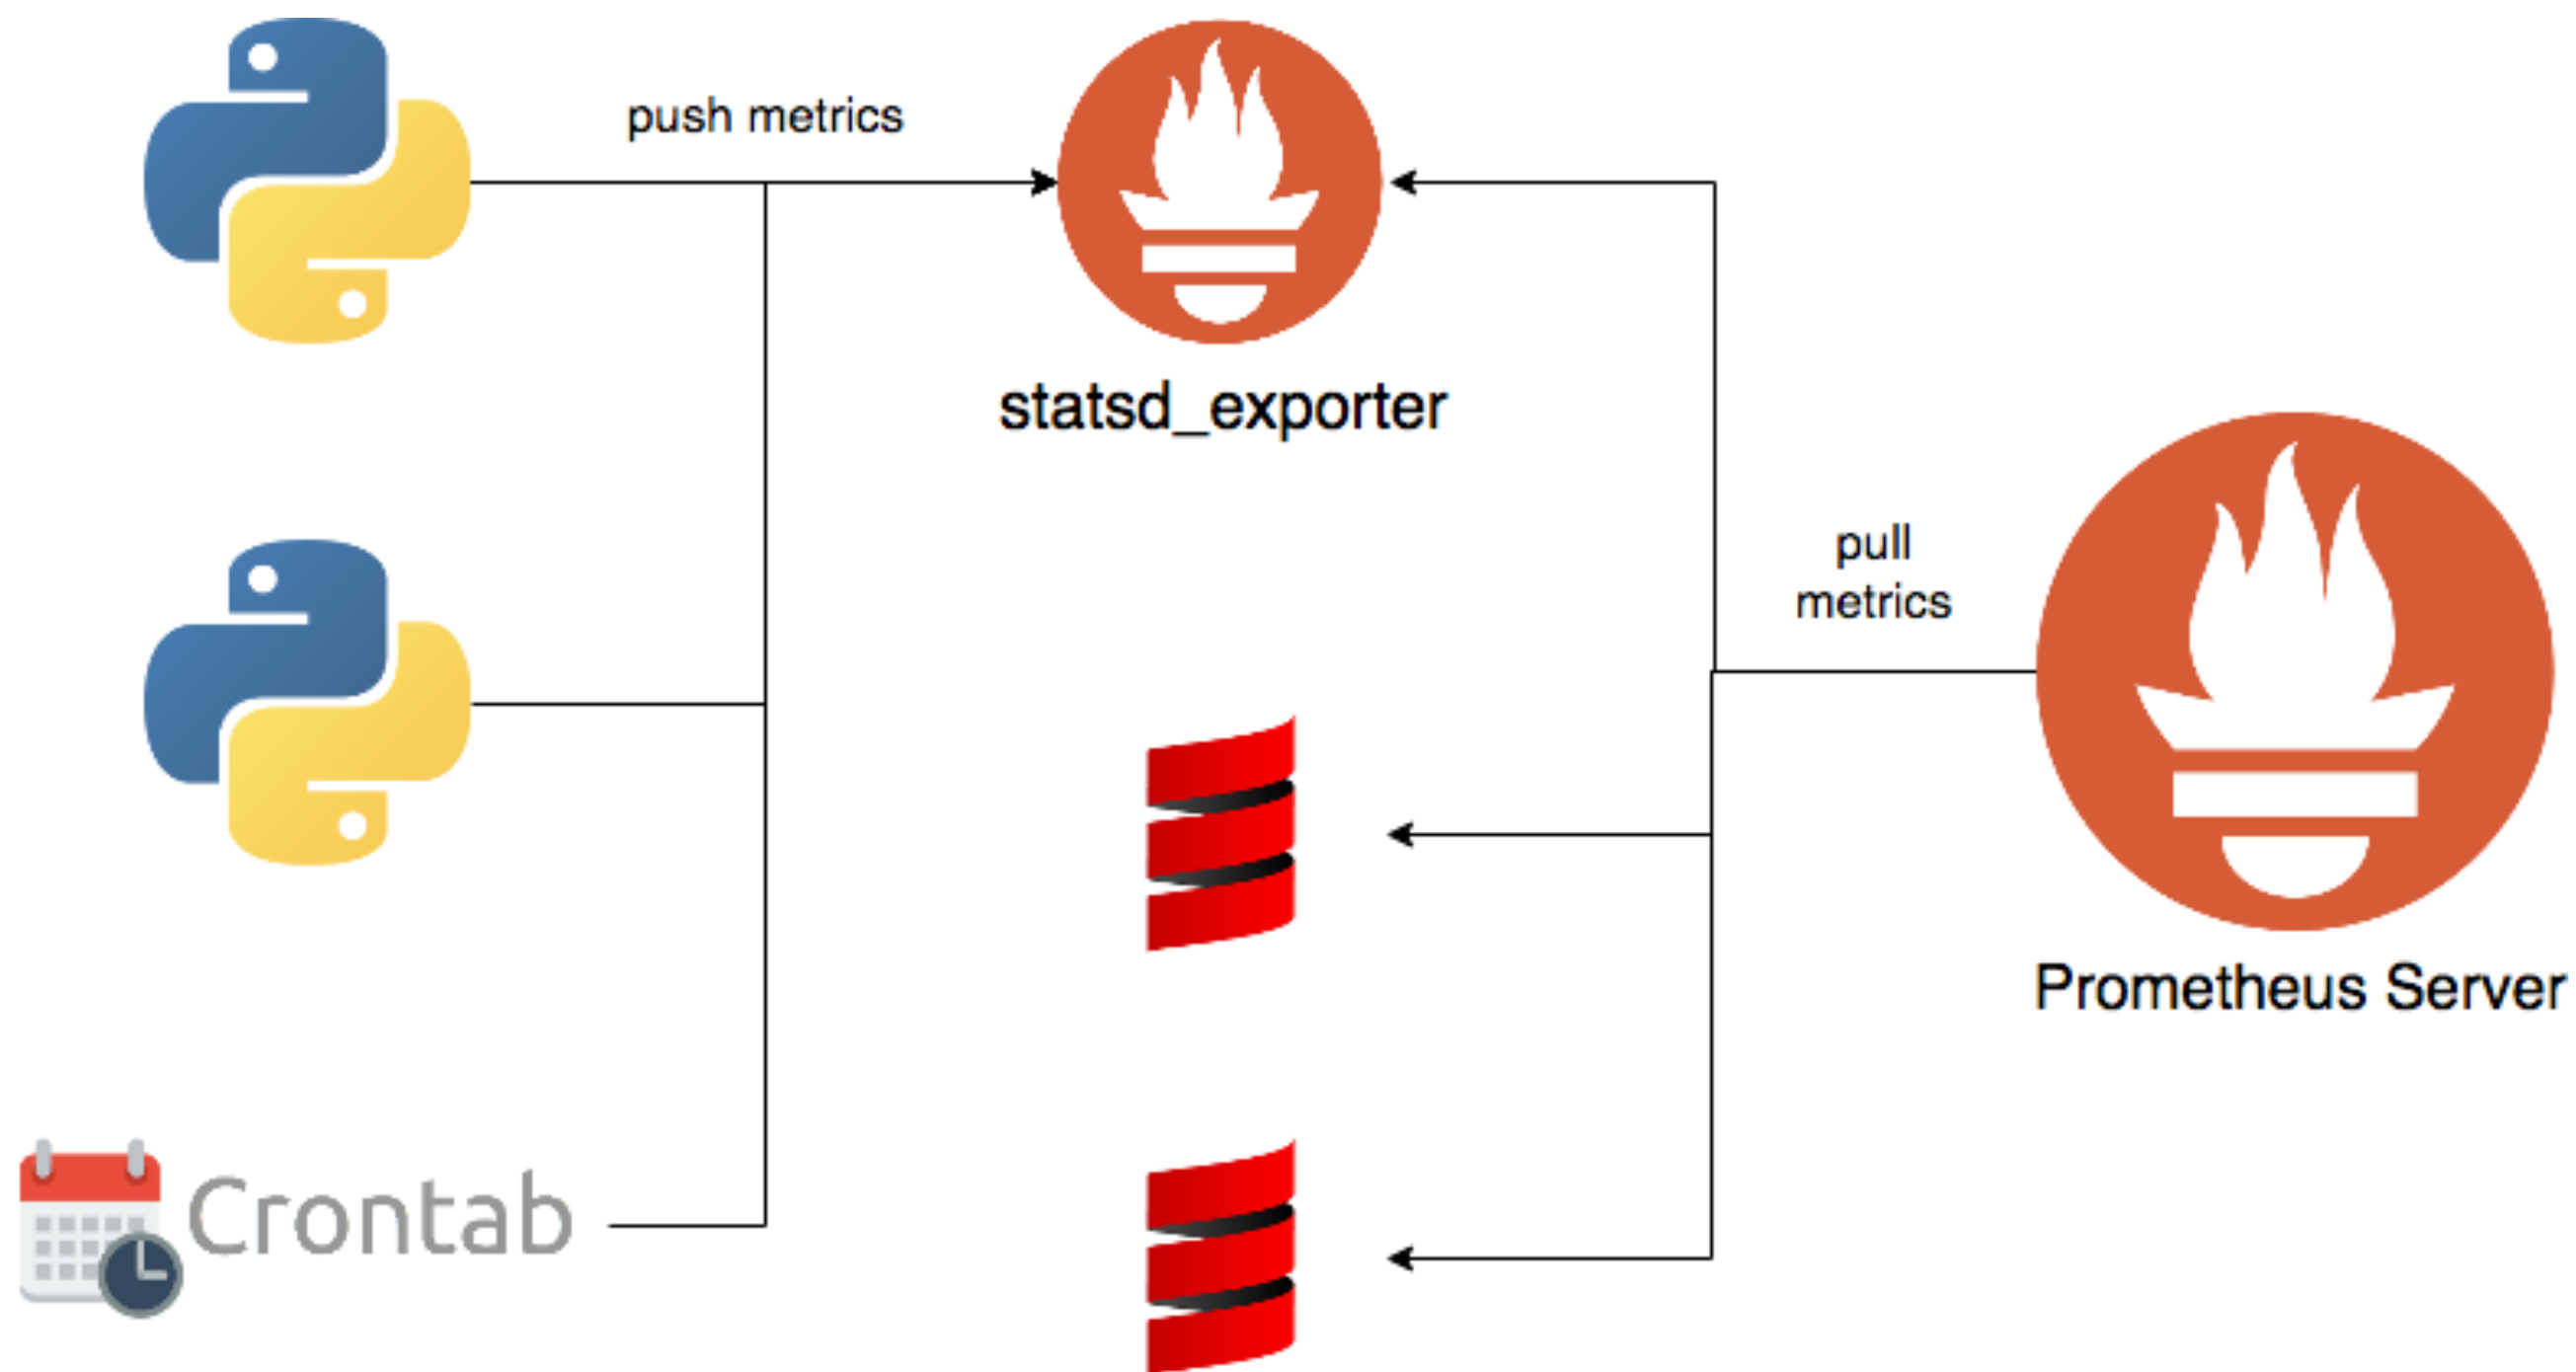
```
$ docker run -d \
  -p 9102:9102 \
  -p 9125:9125 \
  -p 9125:9125/udp \
  prom/statsd-exporter
```

```
1 import aiologstatsd
2 from aiohttp import web
3
4
5 async def fire(request):
6     request.app["statsd"].increment(
7         "reqs_total", tags={"view": "fire"}
8     )
9     return web.Response(body=b"Fire!")
10
11
12 async def statsd_client(app):
13     app["statsd"] = aiologstatsd.Client(host="0.0.0.0", port=9125)
14     await app["statsd"].connect()
15     yield
16     await app["statsd"].close()
17
18
19 async def get_application():
20     app = web.Application()
21     app.add_routes([web.get("/fire", fire)])
22     app.cleanup_ctx.append(statsd_client)
23     return app
```

```
$ gunicorn app_statsd:get_application \  
    --workers=4 \  
    --worker-class=aiohttp.GunicornWebWorker
```

```
$ for i in `seq 30`; do http :8000/fire; done
```

```
$ http :9102/metrics | grep reqs_total  
# HELP reqs_total Metric autogenerated by statsd_exporter.  
# TYPE reqs_total counter  
reqs_total{view="fire"} 30
```



StatsD gauge	→	Prometheus gauge
StatsD counter	→	Prometheus counter
StatsD timer	→	Prometheus summary

StatsD gauge	→	Prometheus gauge
StatsD counter	→	Prometheus counter
StatsD timer	→	Prometheus histogram



# histogram vs summary

- <https://prometheus.io/docs/practices/histograms/>
- <https://povilasv.me/prometheus-tracking-request-duration/>

defaults:

timer\_type: histogram

buckets: [.05, .1, .25, .5, 1, 2.5]

Bonus part!

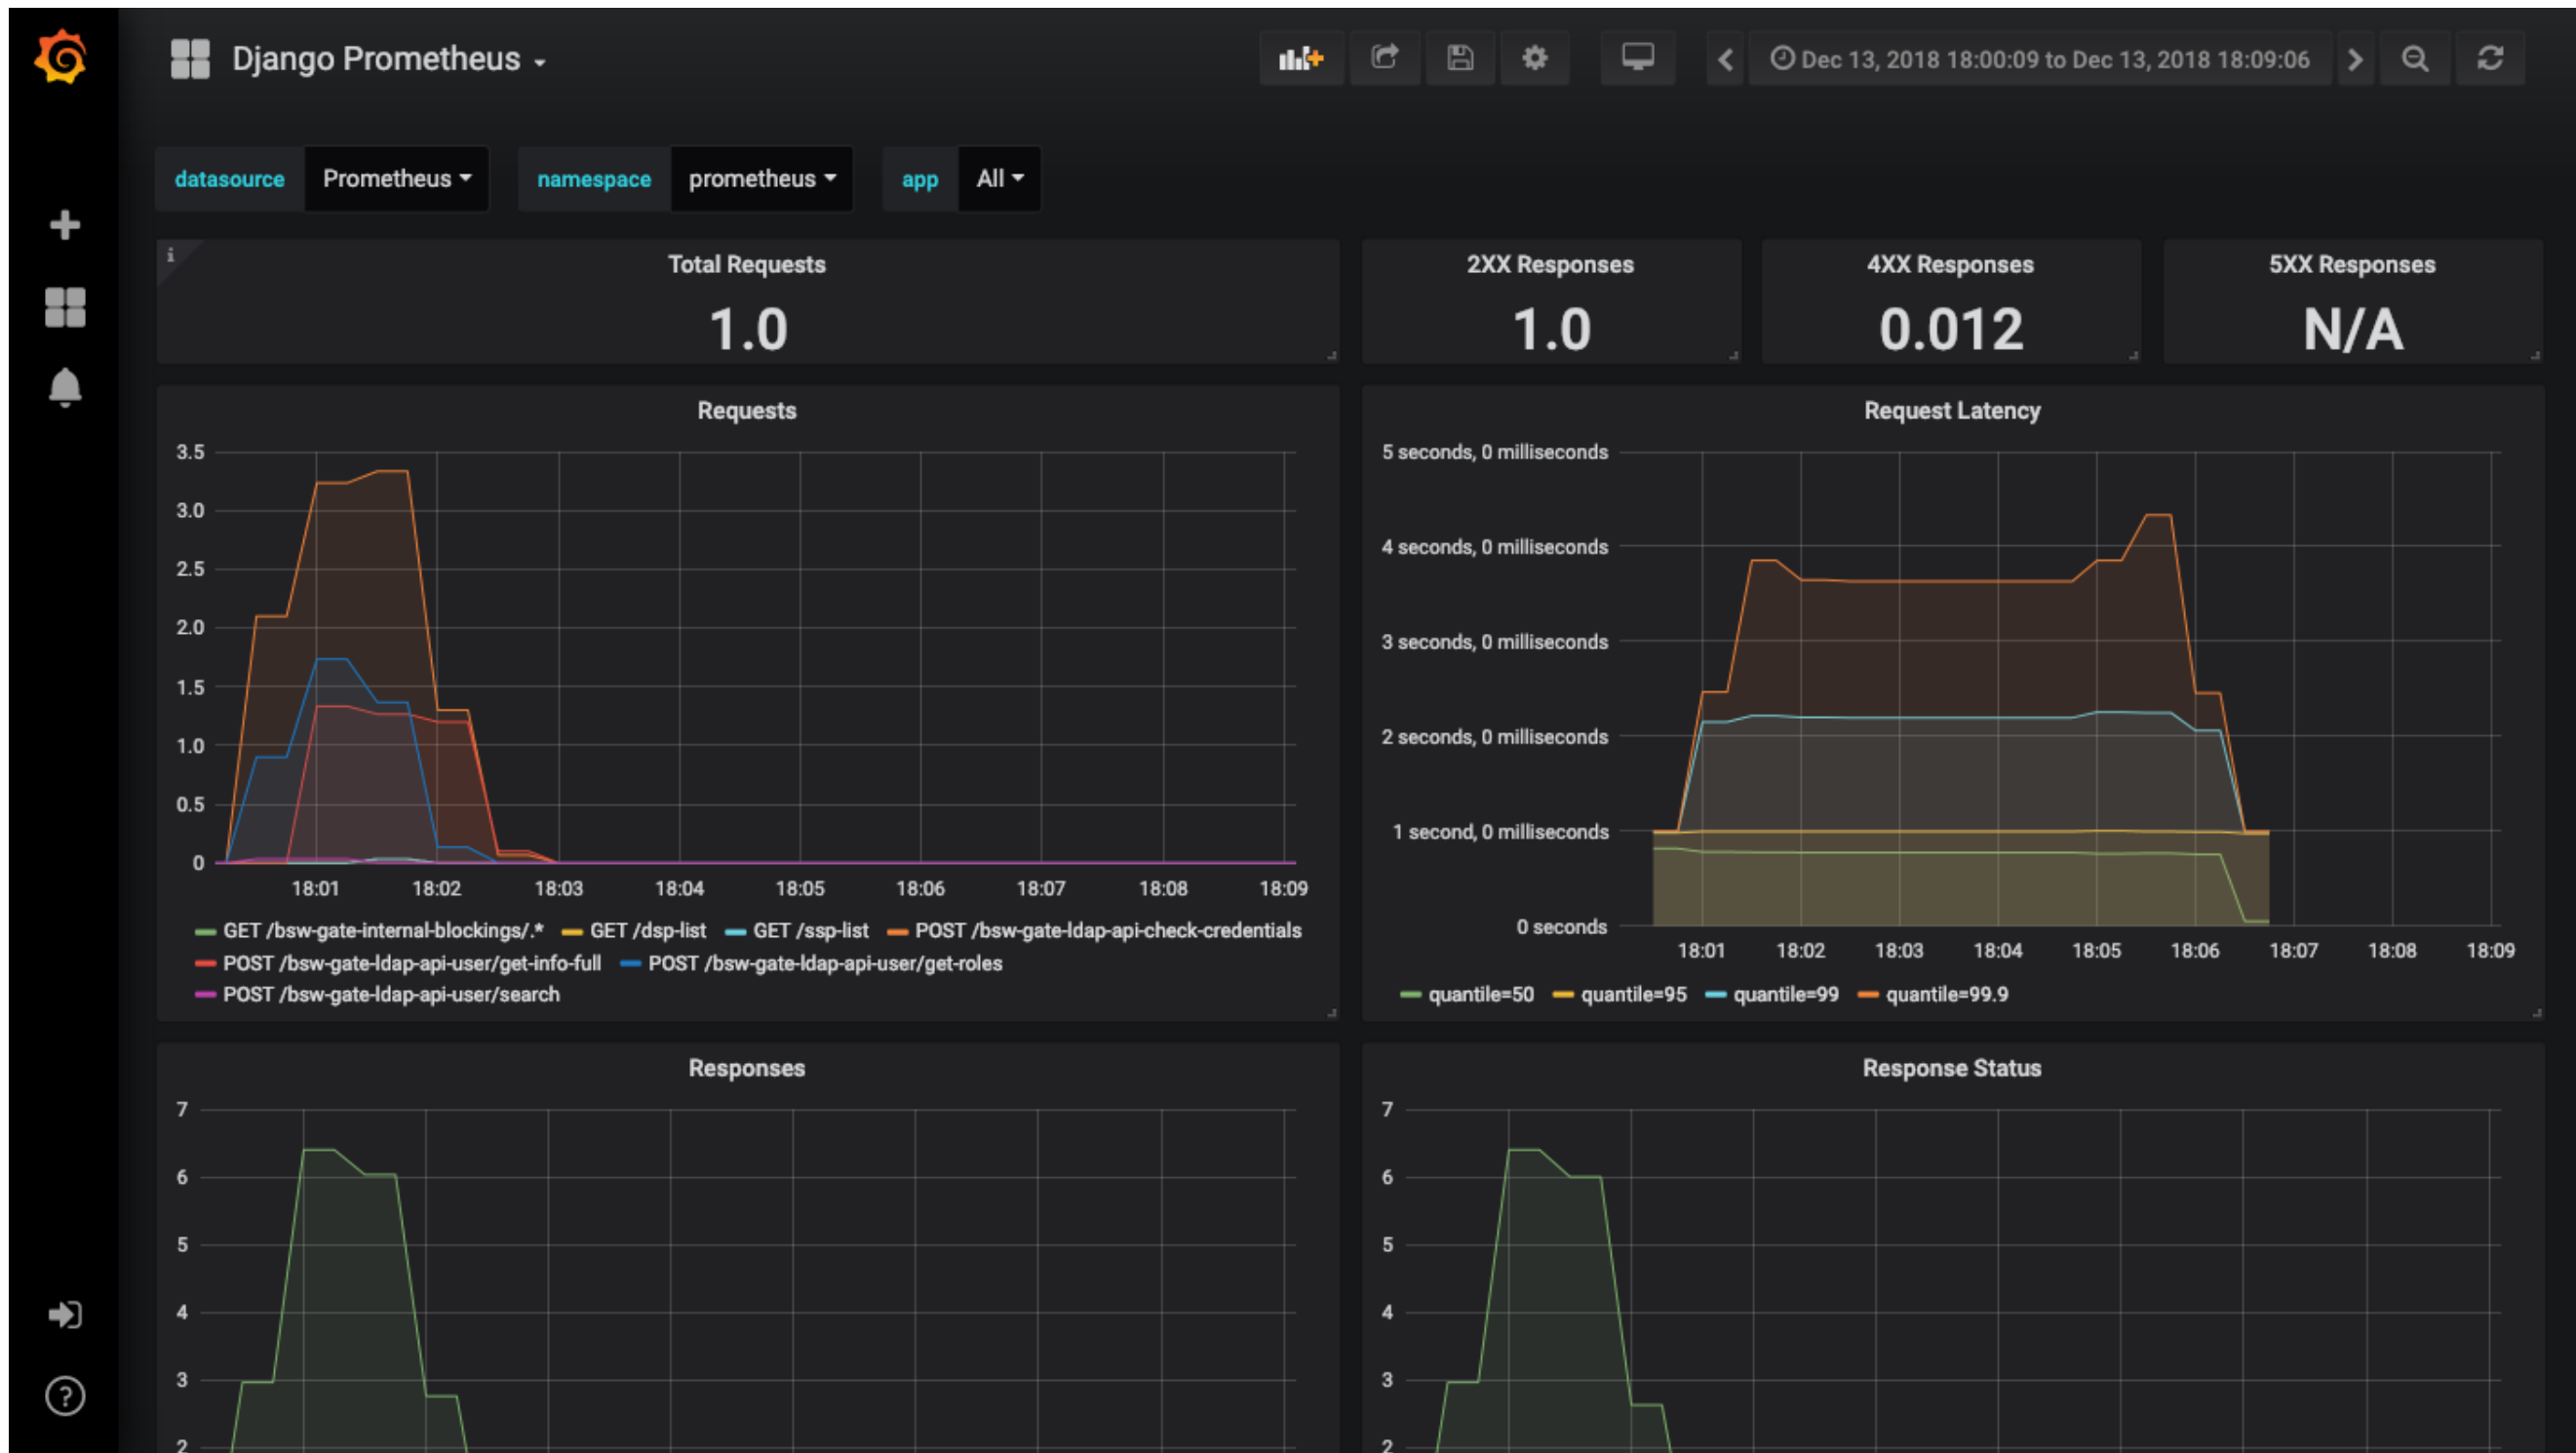
# Gunicorn

- <https://medium.com/@damianmyerscough/monitoring-gunicorn-with-prometheus-789954150069>
- <https://docs.gunicorn.org/en/stable/instrumentation.html>
- `gunicorn --statsd-host=...`
- **AIOHTTP not supported...**
- **Labels not supported...**

# Sentry

- <https://docs.sentry.io/server/internal-metrics/>
- Sentry provides an abstraction called 'metrics' which is used for internal monitoring, generally timings and various counters.

# Grafana



# Grafana

- [https://grafana.com/dashboards?  
dataSource=prometheus](https://grafana.com/dashboards?dataSource=prometheus)
- ~5-10 minutes!

Without StatsD?



# Prometheus + Consul

- <http://blog.archer.onl/article/collecting-prometheus-metrics-for-python-services-behind-gunicorn-using-consul/>
- <https://www.consul.io/>
- [https://github.com/hynek/prometheus\\_async](https://github.com/hynek/prometheus_async)

<https://github.com/Gr1N/talks>

# Questions?

<https://github.com/Gr1N>